

Docker Fundamentals - Comprehensive Notes

1. Introduction to Docker

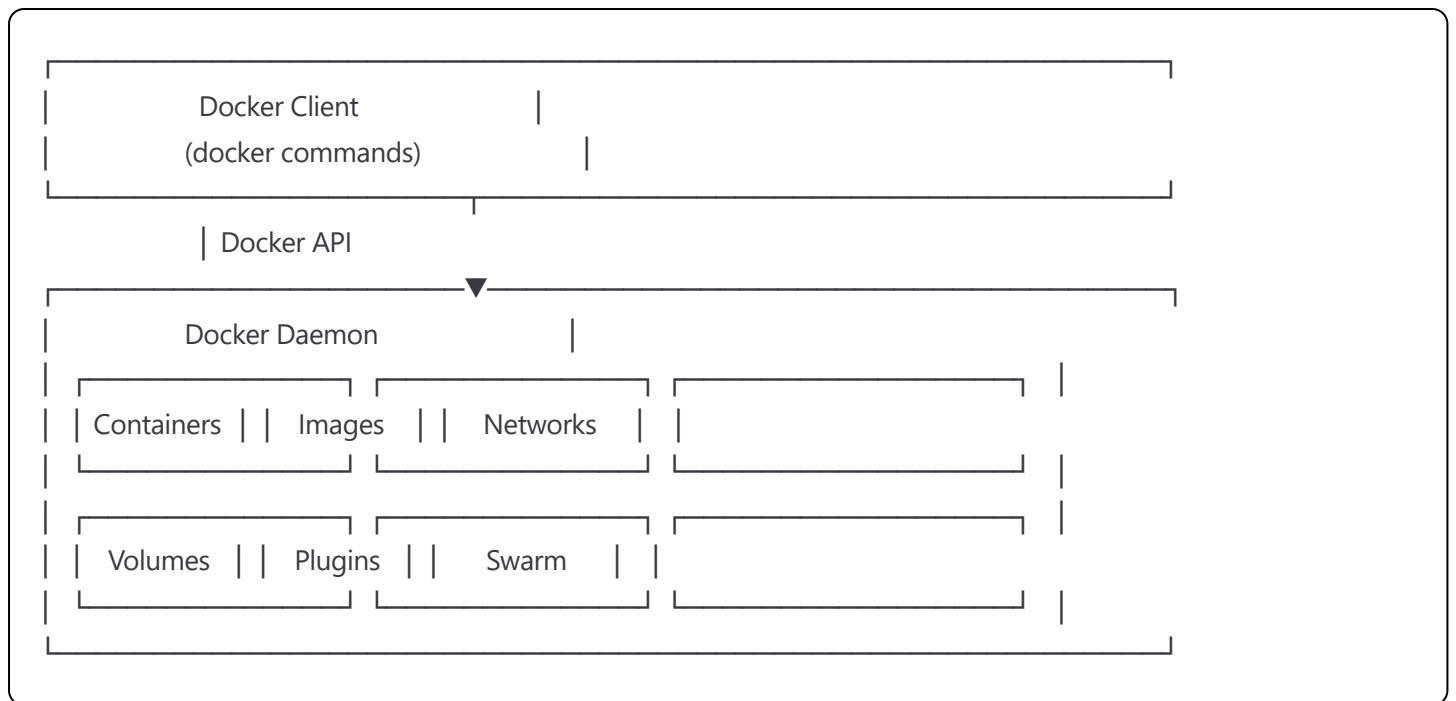
What is Docker?

- **Docker:** A containerization platform that packages applications and their dependencies into lightweight, portable containers
- **Container:** A standardized unit of software that packages code, runtime, system tools, libraries, and settings
- **Key Benefits:**
 - Consistency across environments (dev, test, prod)
 - Lightweight compared to virtual machines
 - Fast startup and deployment
 - Scalability and orchestration
 - Version control for applications

Docker vs Virtual Machines

Docker Containers	Virtual Machines
Share host OS kernel	Each VM has full OS
Lightweight (MBs)	Heavy (GBs)
Fast startup (seconds)	Slow startup (minutes)
Less resource usage	More resource usage
Process-level isolation	Hardware-level isolation

Docker Architecture



2. Core Docker Concepts

Images

- **Definition:** Read-only template used to create containers
- **Layered Architecture:** Built in layers using Union File System
- **Base Image:** Starting point (e.g., ubuntu, alpine, python)
- **Image Registry:** Repository for storing and sharing images (Docker Hub, ECR, etc.)

Containers

- **Definition:** Running instance of an image
- **Lifecycle:** Created → Running → Stopped → Removed
- **Isolation:** Process, network, and file system isolation
- **Stateless:** Containers should be stateless; data persisted via volumes

Dockerfile

- **Definition:** Text file containing instructions to build an image
- **Declarative:** Describes the desired state of the image
- **Caching:** Docker caches layers to speed up builds
- **Best Practices:** Minimize layers, use specific tags, clean up in same layer

Docker Registry

- **Docker Hub:** Default public registry
- **Private Registries:** ECR, GCR, Harbor, etc.
- **Image Naming:** `registry/namespace/repository:tag`

3. Docker Installation and Setup

Installation Steps (Ubuntu/Debian)

```
bash

# Remove old versions
sudo apt remove docker docker-engine docker.io containerd runc

# Update package index
sudo apt update

# Install packages for HTTPS repository
sudo apt install ca-certificates curl gnupg lsb-release

# Add Docker GPG key
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /usr/share/keyrings/docker-archive-keyring.gpg

# Add Docker repository
echo "deb [arch=$(dpkg --print-architecture) signed-by=/usr/share/keyrings/docker-archive-keyring.gpg] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable" | sudo tee /etc/apt/sources.list.d/docker.list >/dev/null

# Install Docker Engine
sudo apt update
sudo apt install docker-ce docker-ce-cli containerd.io docker-compose-plugin

# Start and enable Docker
sudo systemctl start docker
sudo systemctl enable docker

# Add user to docker group (optional)
sudo usermod -aG docker $USER
```

Verify Installation

```
bash

# Check Docker version
docker --version

# Run hello-world container
docker run hello-world

# Check Docker info
docker info
```

4. Essential Docker Commands

Image Management

```
bash

# Search for images
docker search ubuntu

# Pull image from registry
docker pull ubuntu:20.04
docker pull nginx:latest

# List local images
docker images
docker images -a      # Show all images including intermediates

# Remove images
docker rmi image_name:tag
docker rmi image_id
docker image prune     # Remove unused images
docker image prune -a  # Remove all unused images

# Build image from Dockerfile
docker build -t my-app:1.0 .
docker build -t my-app:1.0 -f Dockerfile.prod .

# Tag image
docker tag source_image:tag target_image:tag

# Push image to registry
docker push my-app:1.0

# Image history
docker history image_name

# Inspect image
docker inspect image_name
```

Container Management

```
bash
```

Run container

```
docker run ubuntu:20.04      # Run and exit
docker run -it ubuntu:20.04  # Interactive mode
docker run -d nginx:latest    # Detached mode (background)
docker run --name my-container nginx # Named container
docker run -p 8080:80 nginx   # Port mapping
docker run -v /host/path:/container/path ubuntu # Volume mount
```

List containers

```
docker ps      # Running containers
docker ps -a   # All containers
docker ps -q   # Only container IDs
```

Start/Stop containers

```
docker start container_name
docker stop container_name
docker restart container_name
docker pause container_name
docker unpause container_name
```

Remove containers

```
docker rm container_name
docker rm -f container_name # Force remove running container
docker container prune      # Remove all stopped containers
```

Execute commands in running container

```
docker exec -it container_name /bin/bash
docker exec container_name ls -la
docker exec -u root container_name /bin/bash
```

View container logs

```
docker logs container_name
docker logs -f container_name # Follow logs
docker logs --tail 100 container_name
```

Inspect container

```
docker inspect container_name
```

Container statistics

```
docker stats
docker stats container_name
```

Copy files to/from container

```
docker cp file.txt container_name:/path/to/destination
docker cp container_name:/path/to/file.txt ./local/path
```

Container Lifecycle Commands

bash

Create container without starting

`docker create --name my-container nginx`

Attach to running container

`docker attach container_name`

Export container as tar

`docker export container_name > container.tar`

Import container from tar

`docker import container.tar new_image:tag`

Commit container changes to new image

`docker commit container_name new_image:tag`

5. Dockerfile Deep Dive

Basic Dockerfile Structure

dockerfile

```
# Use official base image
FROM ubuntu:20.04

# Set maintainer info
LABEL maintainer="your-email@example.com"
LABEL version="1.0"
LABEL description="Sample application"

# Set working directory
WORKDIR /app

# Copy files from host to container
COPY ./app /app
COPY requirements.txt .

# Install dependencies
RUN apt-get update && \
    apt-get install -y python3 python3-pip && \
    pip3 install -r requirements.txt && \
    apt-get clean && \
    rm -rf /var/lib/apt/lists/*

# Set environment variables
ENV APP_ENV=production
ENV PORT=8000

# Expose port
EXPOSE 8000

# Create non-root user
RUN useradd -m -s /bin/bash appuser
USER appuser

# Define volume
VOLUME ["/app/data"]

# Define entry point
ENTRYPOINT ["python3"]
CMD ["app.py"]
```

Dockerfile Instructions

dockerfile

```
# FROM - Base image
FROM python:3.9-slim

# LABEL - Metadata
LABEL version="1.0" maintainer="dev@company.com"

# WORKDIR - Set working directory
WORKDIR /usr/src/app

# COPY - Copy files/directories
COPY . .          # Copy everything
COPY src/ /app/src/  # Copy specific directory

# ADD - Copy and extract archives
ADD app.tar.gz /app/

# RUN - Execute commands during build
RUN pip install --no-cache-dir -r requirements.txt

# ENV - Set environment variables
ENV PYTHONPATH=/usr/src/app
ENV DEBUG=False

# EXPOSE - Document port usage
EXPOSE 8000

# USER - Set user context
USER nobody

# VOLUME - Create mount point
VOLUME ["/data", "/logs"]

# ENTRYPOINT - Configure container executable
ENTRYPOINT ["python"]

# CMD - Default command/arguments
CMD ["app.py"]

# ARG - Build-time variables
ARG BUILD_DATE
ARG VERSION=1.0
```

Multi-stage Dockerfile Example

dockerfile


```
# Build stage
```

```
FROM node:16-alpine AS builder
```

```
WORKDIR /app
```

```
COPY package*.json ./
```

```
RUN npm ci --only=production
```

```
# Production stage
```

```
FROM node:16-alpine AS production
```

```
WORKDIR /app
```

```
COPY --from=builder /app/node_modules ./node_modules
```

```
COPY ..
```

```
EXPOSE 3000
```

```
CMD ["node", "server.js"]
```

6. Docker Networking

Network Types

```
bash
```

```
# List networks
```

```
docker network ls
```

```
# Default networks:
```

```
# bridge - Default network for containers
```

```
# host - Use host networking
```

```
# none - No networking
```

Network Management

```
bash
```

Create custom network

`docker network create my-network`

`docker network create --driver bridge my-bridge-network`

Run container with custom network

`docker run -d --name web --network my-network nginx`

Connect container to network

`docker network connect my-network container_name`

Disconnect from network

`docker network disconnect my-network container_name`

Inspect network

`docker network inspect my-network`

Remove network

`docker network rm my-network`

Port Mapping

bash

Map container port to host port

`docker run -p 8080:80 nginx` *# Host:Container*

`docker run -p 127.0.0.1:8080:80 nginx` *# Bind to specific interface*

`docker run -P nginx` *# Map all exposed ports randomly*

7. Docker Volumes

Volume Types

1. **Named Volumes:** Managed by Docker
2. **Bind Mounts:** Direct host directory mapping
3. **tmpfs Mounts:** Temporary in-memory storage

Volume Management

bash

Create named volume

`docker volume create my-volume`

List volumes

`docker volume ls`

Inspect volume

`docker volume inspect my-volume`

Remove volume

`docker volume rm my-volume`

`docker volume prune` *# Remove unused volumes*

Use named volume

`docker run -v my-volume:/app/data nginx`

Use bind mount

`docker run -v /host/path:/container/path nginx`

`docker run -v $(pwd):/app nginx` *# Current directory*

Use tmpfs mount

`docker run --tmpfs /tmp nginx`

Volume Examples

bash

Database with persistent storage

```
docker run -d \  
  --name postgres-db \  
  -v postgres-data:/var/lib/postgresql/data \  
  -e POSTGRES_PASSWORD=mypassword \  
  postgres:13
```

Development environment with code binding

```
docker run -d \  
  --name dev-container \  
  -v $(pwd)/src:/app/src \  
  -v node_modules:/app/node_modules \  
  -p 3000:3000 \  
  node:16-alpine
```

8. Docker Compose

What is Docker Compose?

- **Definition:** Tool for defining and running multi-container Docker applications
- **YAML Configuration:** Use `docker-compose.yml` file
- **Service Management:** Manage multiple services as a single application
- **Networking:** Automatic network creation for services

Basic docker-compose.yml

yaml

version: '3.8'

services:

web:

build: .

ports:

- "8000:8000"

depends_on:

- db

- redis

environment:

- DATABASE_URL=postgresql://user:pass@db:5432/mydb

volumes:

- ./app:/app

- static_volume:/app/static

db:

image: postgres:13

environment:

POSTGRES_DB: mydb

POSTGRES_USER: user

POSTGRES_PASSWORD: pass

volumes:

- postgres_data:/var/lib/postgresql/data

redis:

image: redis:6-alpine

ports:

- "6379:6379"

volumes:

postgres_data:

static_volume:

networks:

default:

driver: bridge

Docker Compose Commands

bash

```
# Start services
docker-compose up
docker-compose up -d      # Detached mode
docker-compose up --build  # Rebuild images

# Stop services
docker-compose down
docker-compose down -v     # Remove volumes
docker-compose down --rmi all # Remove images

# View services
docker-compose ps
docker-compose logs
docker-compose logs -f web  # Follow logs for specific service

# Execute commands
docker-compose exec web bash
docker-compose exec db psql -U user -d mydb

# Scale services
docker-compose up -d --scale web=3

# Build images
docker-compose build
docker-compose build --no-cache

# Pull images
docker-compose pull
```

9. Docker Best Practices

Dockerfile Best Practices

```
dockerfile
```

Use specific tags, not 'latest'

FROM python:3.9-slim

Use multi-stage builds

FROM node:16-alpine **AS** builder

... build steps ...

FROM node:16-alpine **AS** production

COPY --from=builder /app/dist ./dist

Minimize layers

RUN apt-get update && \

apt-get install -y package1 package2 && \

apt-get clean && \

rm -rf /var/lib/apt/lists/*

Don't run as root

RUN useradd -m -s /bin/bash appuser

USER appuser

Use .dockerignore

Create .dockerignore file to exclude unnecessary files

Security Best Practices

dockerfile

Use official base images

FROM python:3.9-slim

Don't store secrets in images

Use environment variables or secrets management

Use non-root user

USER 1000:1000

Use specific versions

FROM python:3.9.10-slim

Scan images for vulnerabilities

docker scan image_name

Performance Best Practices

dockerfile

```
# Order layers by change frequency
```

```
FROM python:3.9-slim
```

```
COPY requirements.txt .
```

```
RUN pip install -r requirements.txt
```

```
COPY ..
```

```
# Use .dockerignore
```

```
# Exclude node_modules, .git, etc.
```

```
# Use multi-stage builds for smaller images
```

```
FROM node:16-alpine AS builder
```

```
# ... build steps ...
```

```
FROM nginx:alpine
```

```
COPY --from=builder /app/dist /usr/share/nginx/html
```

10. Container Orchestration Basics

Docker Swarm

```
bash
```

```
# Initialize swarm
```

```
docker swarm init
```

```
# Join swarm as worker
```

```
docker swarm join --token <token> <manager-ip>:2377
```

```
# Deploy stack
```

```
docker stack deploy -c docker-compose.yml my-stack
```

```
# List services
```

```
docker service ls
```

```
# Scale service
```

```
docker service scale my-stack_web=3
```

```
# Remove stack
```

```
docker stack rm my-stack
```

Kubernetes Concepts

- **Pods:** Smallest deployable units
- **Services:** Network access to pods
- **Deployments:** Manage pod replicas
- **ConfigMaps/Secrets:** Configuration management

- **Ingress:** External access to services

11. Docker Registry and Image Management

Docker Hub

```
bash

# Login to Docker Hub
docker login

# Tag and push image
docker tag my-app:latest username/my-app:latest
docker push username/my-app:latest

# Pull specific version
docker pull username/my-app:1.0
```

Private Registry

```
bash

# Run local registry
docker run -d -p 5000:5000 --name registry registry:2

# Tag for private registry
docker tag my-app:latest localhost:5000/my-app:latest

# Push to private registry
docker push localhost:5000/my-app:latest
```

12. Monitoring and Logging

Container Monitoring

```
bash
```

View container resource usage

`docker stats`

View container processes

`docker top` container_name

Monitor container events

`docker events`

Export container metrics

`docker stats --format "table {{.Name}}\t{{.CPUPerc}}\t{{.MemUsage}}"`

Logging

bash

View logs

`docker logs` container_name

Follow logs

`docker logs -f` container_name

Logs with timestamps

`docker logs -t` container_name

Limit log output

`docker logs --tail 50` container_name

`docker logs --since 2h` container_name

13. Troubleshooting Common Issues

Container Won't Start

bash

Check container logs

`docker logs container_name`

Check container configuration

`docker inspect container_name`

Run container interactively

`docker run -it image_name /bin/bash`

Check if port is already in use

`netstat -tulpn | grep :8080`

Image Build Issues

bash

Build with no cache

`docker build --no-cache -t my-app .`

Build with verbose output

`docker build --progress=plain -t my-app .`

Check build context

`docker build -t my-app --file Dockerfile.debug .`

Storage Issues

bash

Check disk usage

`docker system df`

Clean up unused resources

`docker system prune`

`docker system prune -a` *# Remove all unused images*

Remove specific resources

`docker container prune`

`docker image prune`

`docker volume prune`

`docker network prune`

14. Common Docker Commands Cheat Sheet

Quick Reference

bash

Images

`docker images` *# List images*
`docker pull image:tag` *# Pull image*
`docker rmi image:tag` *# Remove image*
`docker build -t name:tag .` *# Build image*

Containers

`docker ps` *# List running containers*
`docker ps -a` *# List all containers*
`docker run -d image:tag` *# Run container detached*
`docker stop container_name` *# Stop container*
`docker rm container_name` *# Remove container*
`docker exec -it container bash` *# Execute command in container*

System

`docker info` *# Docker system info*
`docker version` *# Docker version*
`docker system prune` *# Clean up unused resources*

Logs and monitoring

`docker logs container_name` *# View logs*
`docker stats` *# View resource usage*
`docker top container_name` *# View processes*

15. Interview Questions and Answers

Basic Questions

1. What is Docker and why use it?

- Containerization platform for packaging applications
- Benefits: Consistency, portability, scalability, isolation

2. Difference between container and image?

- Image: Read-only template
- Container: Running instance of an image

3. What is Dockerfile?

- Text file with instructions to build Docker image
- Contains commands like FROM, RUN, COPY, CMD

Intermediate Questions

1. Explain Docker architecture

- Client-server architecture

- Docker daemon manages containers, images, networks
- Docker client communicates via REST API

2. What are Docker volumes?

- Persistent storage mechanism
- Types: Named volumes, bind mounts, tmpfs mounts

3. How does Docker networking work?

- Bridge network by default
- Custom networks for isolation
- Port mapping for external access

Advanced Questions

1. Multi-stage Docker builds

- Multiple FROM instructions in Dockerfile
- Reduces final image size
- Separates build and runtime environments

2. Docker security best practices

- Use official base images
- Don't run as root
- Scan images for vulnerabilities
- Use secrets management

3. Container orchestration

- Docker Swarm for clustering
- Kubernetes for advanced orchestration
- Service discovery and load balancing

16. Practical Examples

Python Web Application

dockerfile

```
FROM python:3.9-slim
```

```
WORKDIR /app
```

```
COPY requirements.txt .
```

```
RUN pip install --no-cache-dir -r requirements.txt
```

```
COPY . .
```

```
RUN useradd -m -s /bin/bash appuser
```

```
USER appuser
```

```
EXPOSE 8000
```

```
CMD ["python", "app.py"]
```

Node.js Application with Docker Compose

```
yaml
```

```
version: '3.8'
```

```
services:
```

```
  web:
```

```
    build: .
```

```
    ports:
```

```
      - "3000:3000"
```

```
    environment:
```

```
      - NODE_ENV=production
```

```
      - DATABASE_URL=mongodb://mongo:27017/myapp
```

```
    depends_on:
```

```
      - mongo
```

```
      - redis
```

```
  mongo:
```

```
    image: mongo:4.4
```

```
    volumes:
```

```
      - mongodb_data:/data/db
```

```
  redis:
```

```
    image: redis:6-alpine
```

```
volumes:
```

```
  mongodb_data:
```

NGINX Reverse Proxy

```
dockerfile
```

```
FROM nginx:alpine
```

```
COPY nginx.conf /etc/nginx/nginx.conf
```

```
COPY ./dist /usr/share/nginx/html
```

```
EXPOSE 80
```

```
CMD ["nginx", "-g", "daemon off;"]
```

Summary

Docker is a powerful containerization platform that revolutionizes application deployment and management. Key concepts to master include:

- **Containerization:** Understanding the difference between containers and VMs
- **Images and Containers:** Building, running, and managing Docker images and containers
- **Dockerfile:** Writing efficient and secure Dockerfiles
- **Networking:** Container networking and port mapping
- **Storage:** Volumes and data persistence
- **Orchestration:** Docker Compose for multi-container applications
- **Best Practices:** Security, performance, and operational considerations

Practice these concepts hands-on to build confidence and prepare for technical interviews. Focus on understanding the underlying principles rather than just memorizing commands.